

Amazon Storage Services

S3 – DynamoDB – ElastiCache

Waleed F. Reda

Cloud Storage Services - Introduction

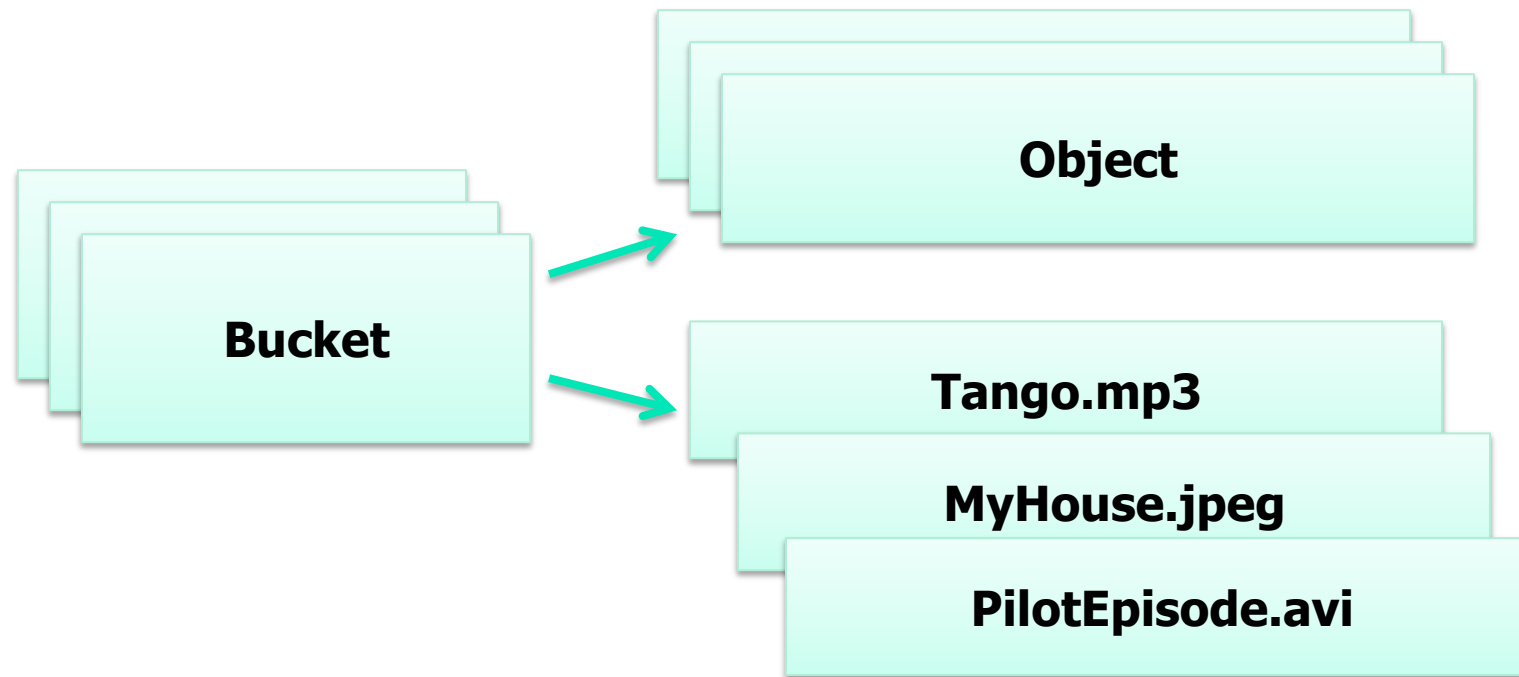
- **Cloud Storage Services** are public infrastructure for the storage of objects
 - Files / blobs / images / videos / etc.
 - Stored in Internet (in a public cloud)
 - Accessible through some API (REST / SDK / etc.)
 - May have front-end for end-user access
 - Could operate as a key-value store
 - Could be free (with limits) or paid (on-demand)

Introduction to Amazon S3

- Amazon S3 == Simple Storage Service
 - On-demand file storage in the AWS cloud
 - Highly-reliable (99.999999999% durability and 99.99% availability)
 - Many APIs: RESTful / SOAP / C# / Java / others
 - Two modes:
 - Normal – more reliable, more expensive
 - Reduced redundancy – cheaper, but less reliable
 - Multiple locations: US, Europe, Asia

Amazon S3 Concepts

- Your cloud storage consists of **buckets**
 - **Objects** are stored in the buckets



Introduction to DynamoDB

- Fully managed NoSQL database service by Amazon
- Database type: Key-value stores
- Designed to address the core problems of database management, performance, scalability, and reliability





NoSQL systems

- Not every data management/analysis problem is best solved *exclusively* using a traditional DBMS
- “NoSQL” = “Not Only SQL”



NoSQL systems

Alternative to traditional relational DBMS

- + Flexible schema
- + Quicker/cheaper to set up
- + Massive scalability
- + Relaxed consistency → higher performance & availability
- No declarative query language → more programming
- Relaxed consistency → fewer guarantees



NoSQL systems

Several incarnations

- Document Stores
- MapReduce framework
- Key-value stores
- Graph database systems

Key-Value Stores


- Extremely simple interface
 - Data model: (key, value) pairs
 - Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)
- Implementation: efficiency, scalability, fault-tolerance
 - Records distributed to nodes based on key
 - Replication
 - Single-record transactions, “eventual consistency”

Key-Value Stores

- Extremely simple interface
 - Data model: (key, value) pairs
 - Operations: Insert(key,value), Fetch(key), Update(key), Delete(key)
- Example systems
 - Google BigTable, Amazon Dynamo, Cassandra, Voldemort, HBase, ...
- Will be covered more thoroughly in upcoming lectures


Features

- Scalable
 - Provisioned Throughput
 - Fully Distributed, Shared Nothing Architecture
- Fast Performance
 - Average service-side latencies < 10 ms
 - The service runs on Solid State Disks - consistent, fast latencies at any scale
- Easy Administration and Cost Effective
 - a fully managed service by Amazon
- Fault-tolerant
 - Synchronous replication across multiple zones in a region




Features

- Flexible
 - Does not have a fixed schema
- Efficient Indexing
 - Every item identified by a primary key
- Strong consistency
 - Implemented with Atomic Counters
 - Disk-only writes



Data model concepts



- Except for the primary key, DynamoDB is schema-less
- Each item can have any number of attributes
- An attribute is a name-value pair
 - can be single valued or multi-valued set

Data model concepts

■ Primary keys

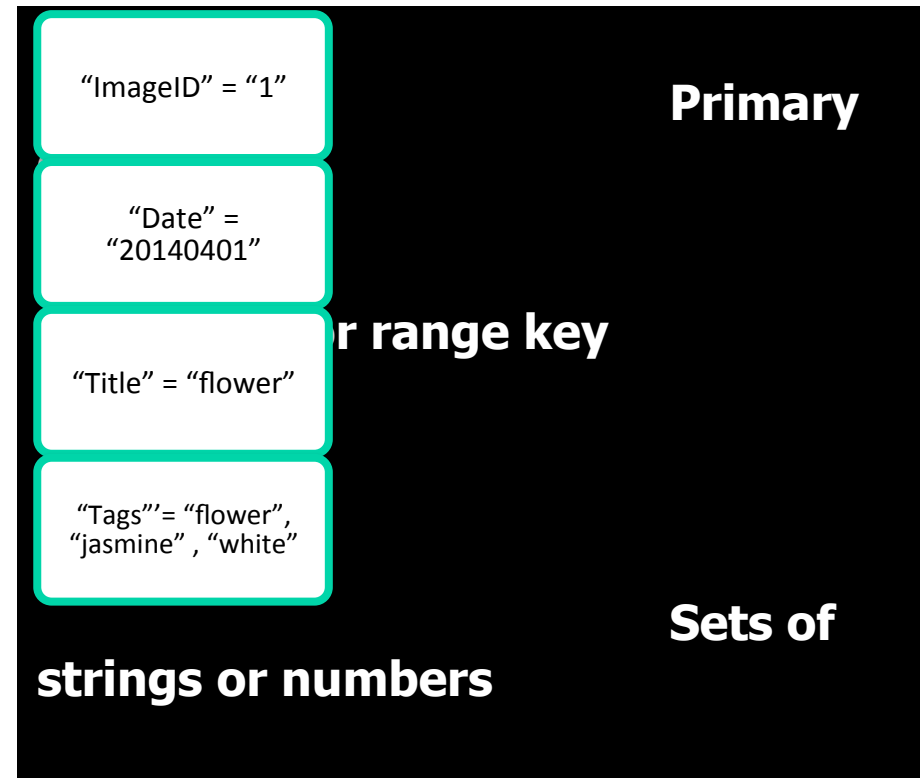
- Hash type primary key
- Hash and range type primary key

■ Secondary Indexes

- Local secondary index
- Global secondary index

■ DynamoDB data types

- Scalar data types
- Multivalued data types



Supported Operations

- Table operations
 - create, update and delete tables
- Item operations
 - add, update and delete items from a table
 - retrieve a single item (GetItem) or multiple items (BatchGetItem)
- Query and Scan
 - query a table using the hash attribute and an optional range filter.
 - If the table has a secondary index, you can also Query the index using its key
 - Scan operation reads every item in the table or secondary index

Supported Operations

- Data Read and Consistency considerations
 - Multiple copies of each item to ensure durability
 - Eventually Consistent Reads
 - Strongly Consistent Reads
- Conditional updates and concurrency control
 - updates made by one client don't overwrite updates made by another client
 - “conditional write” and “atomic counter”



Considerations


- Item size is limited to 64KB
- Attribute values can not be null or empty
- Hash primary key attribute value is limited to 1024 bytes
- Range primary key attribute value is limited to 2048 bytes
- Up to 5 local as well as global secondary indexes per table

Introduction to Amazon ElastiCache

- Memory caching moves repetitive queries to a cache to limit loads on the server.
- The information is in the memory, as opposed to on disk, so it serves up very fast.
- Works on top of Memcached
 - A free open source, high-performance, distributed memory object caching system
- AWS is taking out the complexity by offering Memcached as a service.

Memcached - Introduction

- Very simple
- In-memory key-value object cache service
- “Hot” data from DB stored in cache
- Memory caching tier comprises of a pool of memcached servers each of which is a standalone server that is independent of the other servers in the pool



Memcached – Users and Services



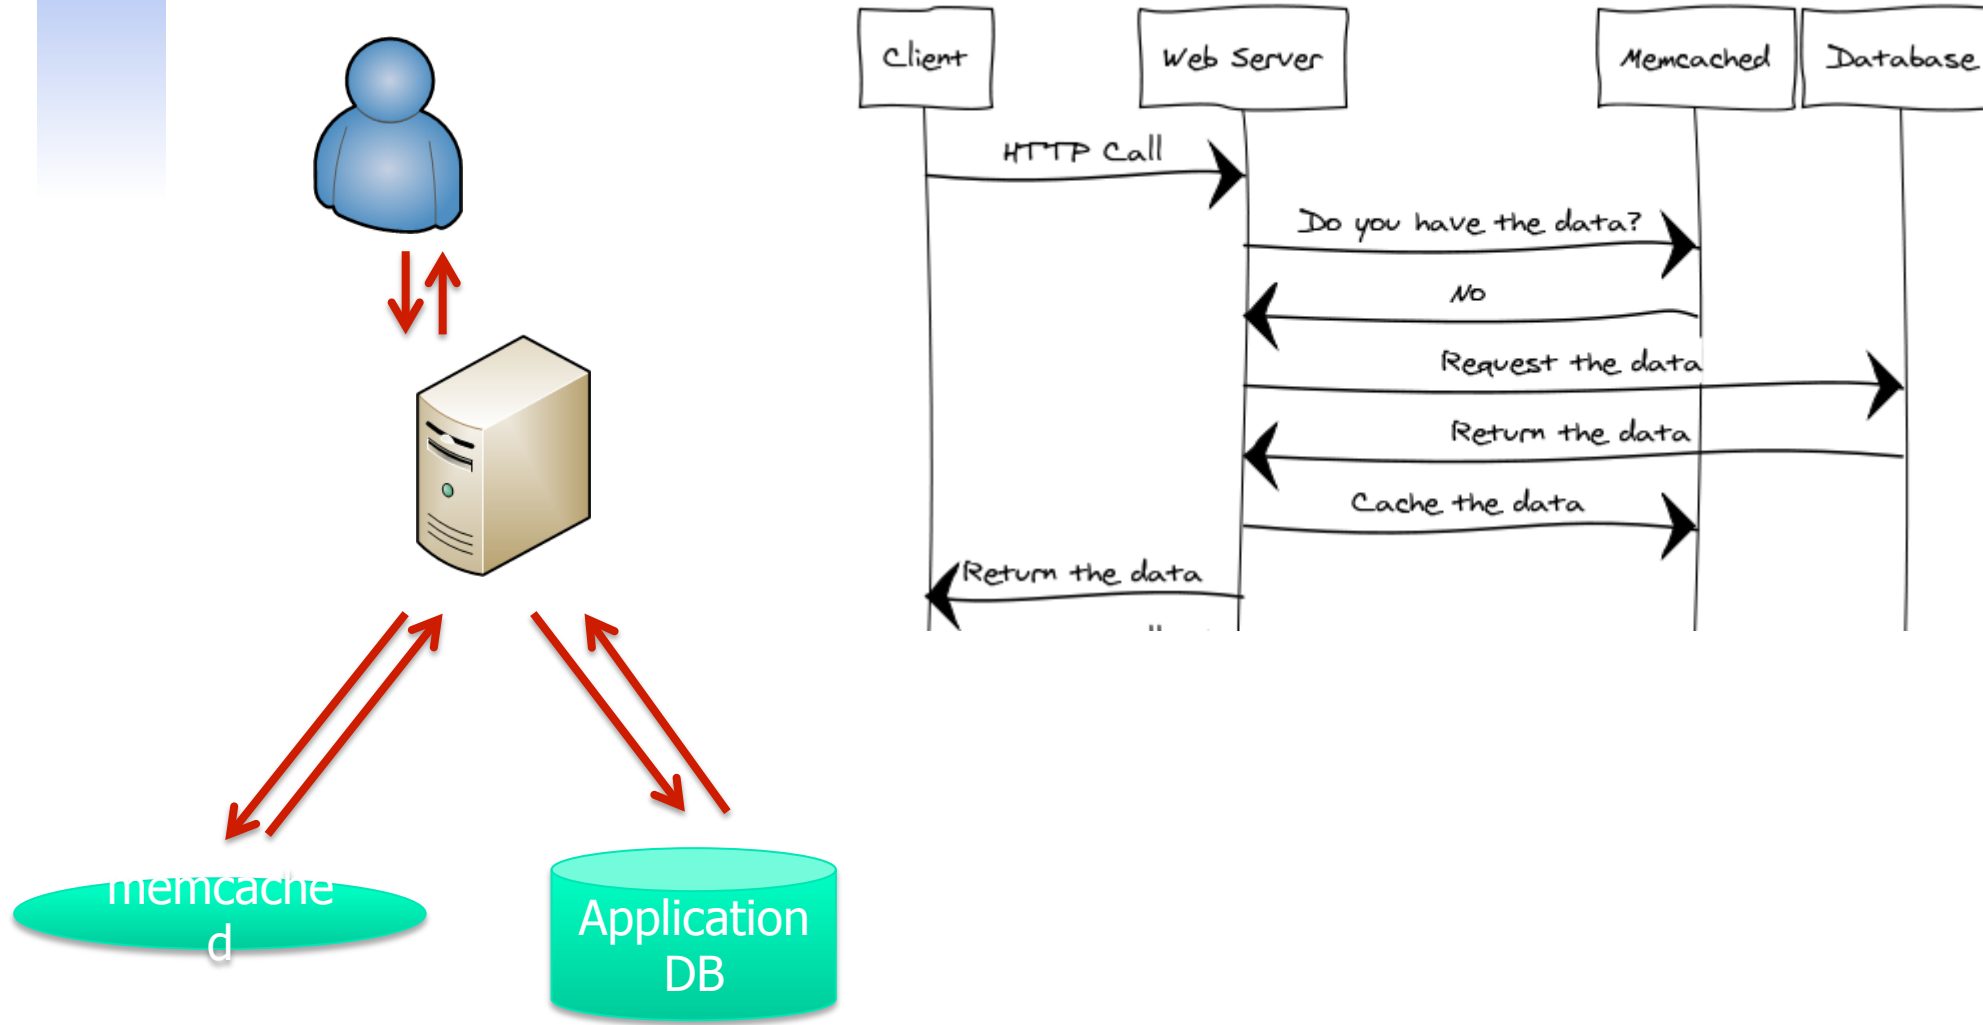
WIKIPEDIA
The Free Encyclopedia



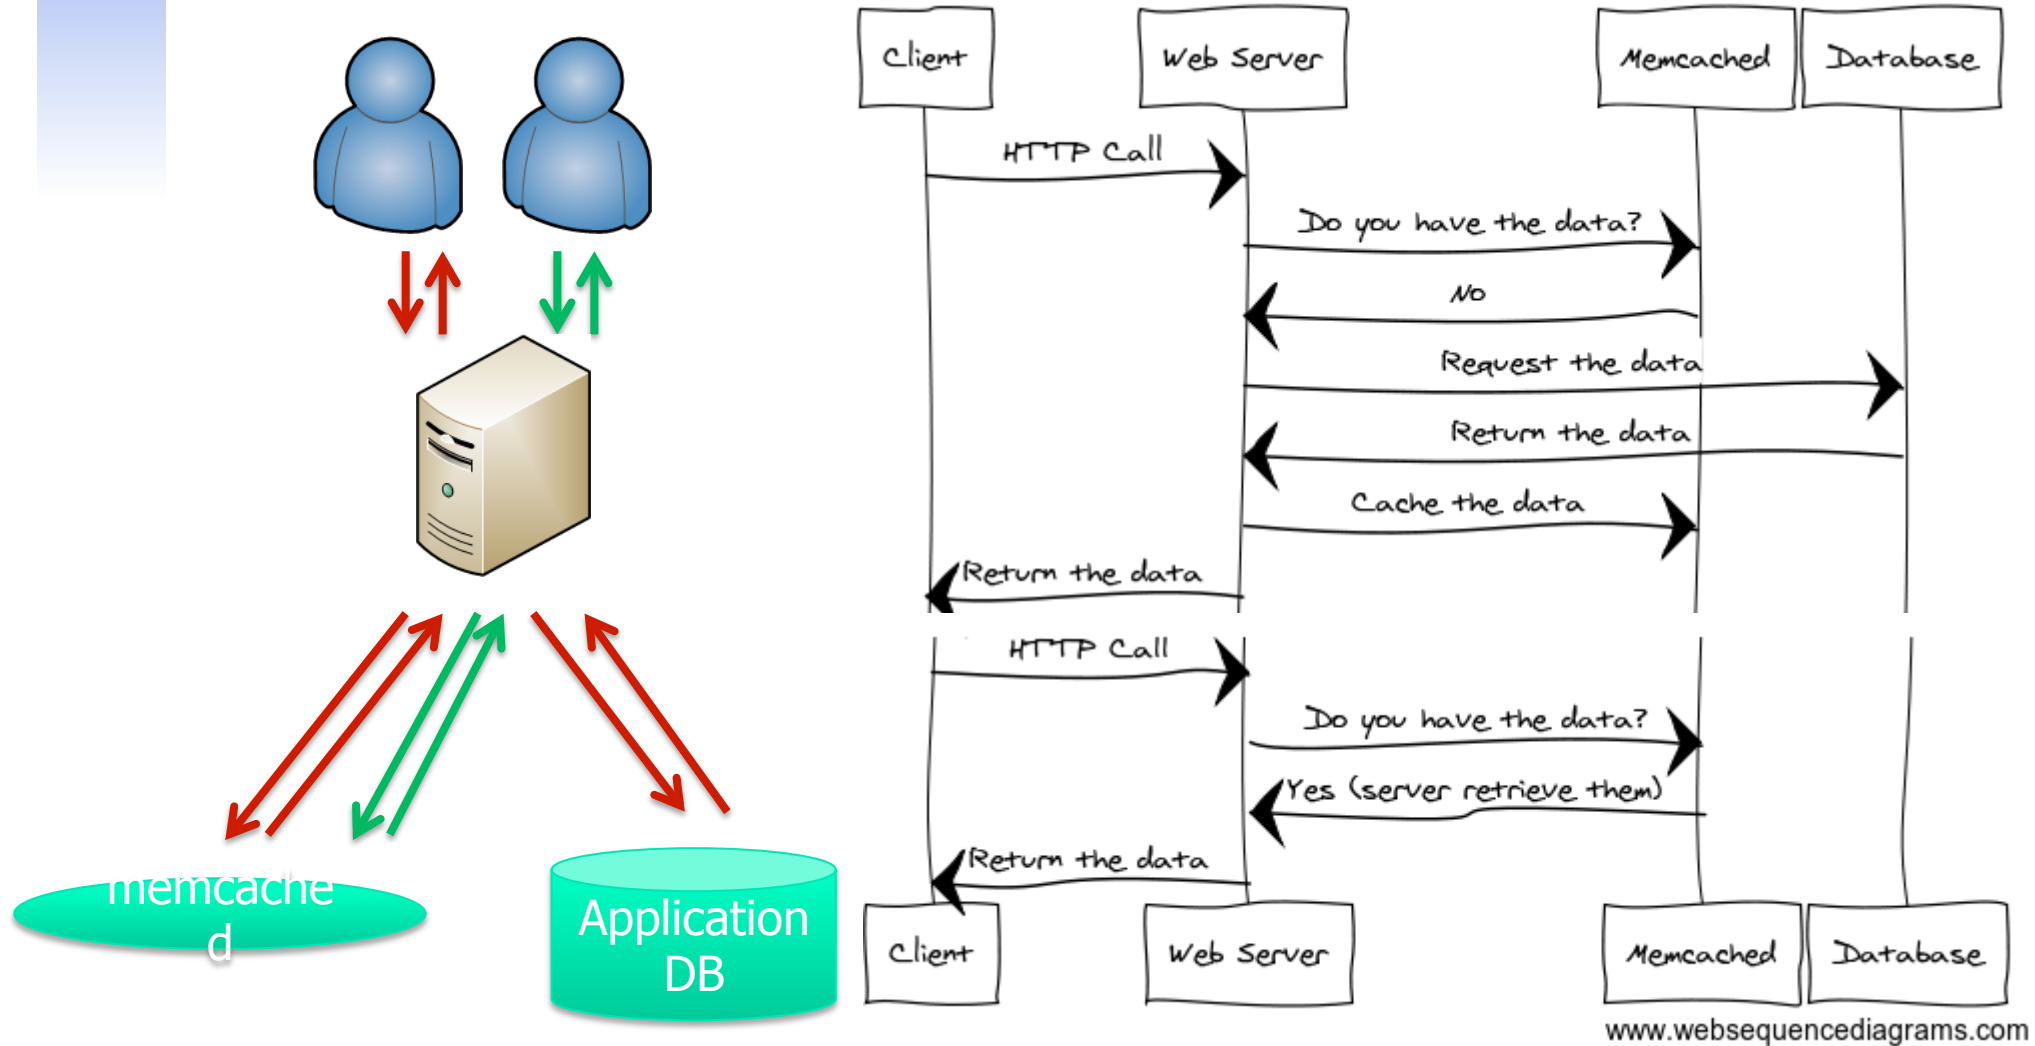
Memcached – Access Characteristics

- Read-mostly workloads
- Writes in order to be durable have to be made to the underlying (non-volatile) data store
- Other studies have shown that writes constitute mere 3-12% of all accesses in large scale key-value stores

Memcached - Example




Memcached - Example



www.websequencediagrams.com

Comparative analysis

Property/ Storage Service	Amazon S3	DynamoDB	ElastiCache
Features	Object storage	Key-value db	Caching of highly utilized data
Limitations	Slower than nosql alternatives	Key/value sizes	Can't be accessed from outside EC2
Use-cases	Cold data	Hot data	Very hot data
Price model			



Refer to Amazon's price calculator: <http://calculator.s3.amazonaws.com/index.html>

Tutorial session

- Please refer to your handouts
- We need to run a total of 3 exercises that spans:
 - S3
 - DynamoDB
 - ElastiCache (Memcached)
- You will be asked to write a few line of code for some added functionality